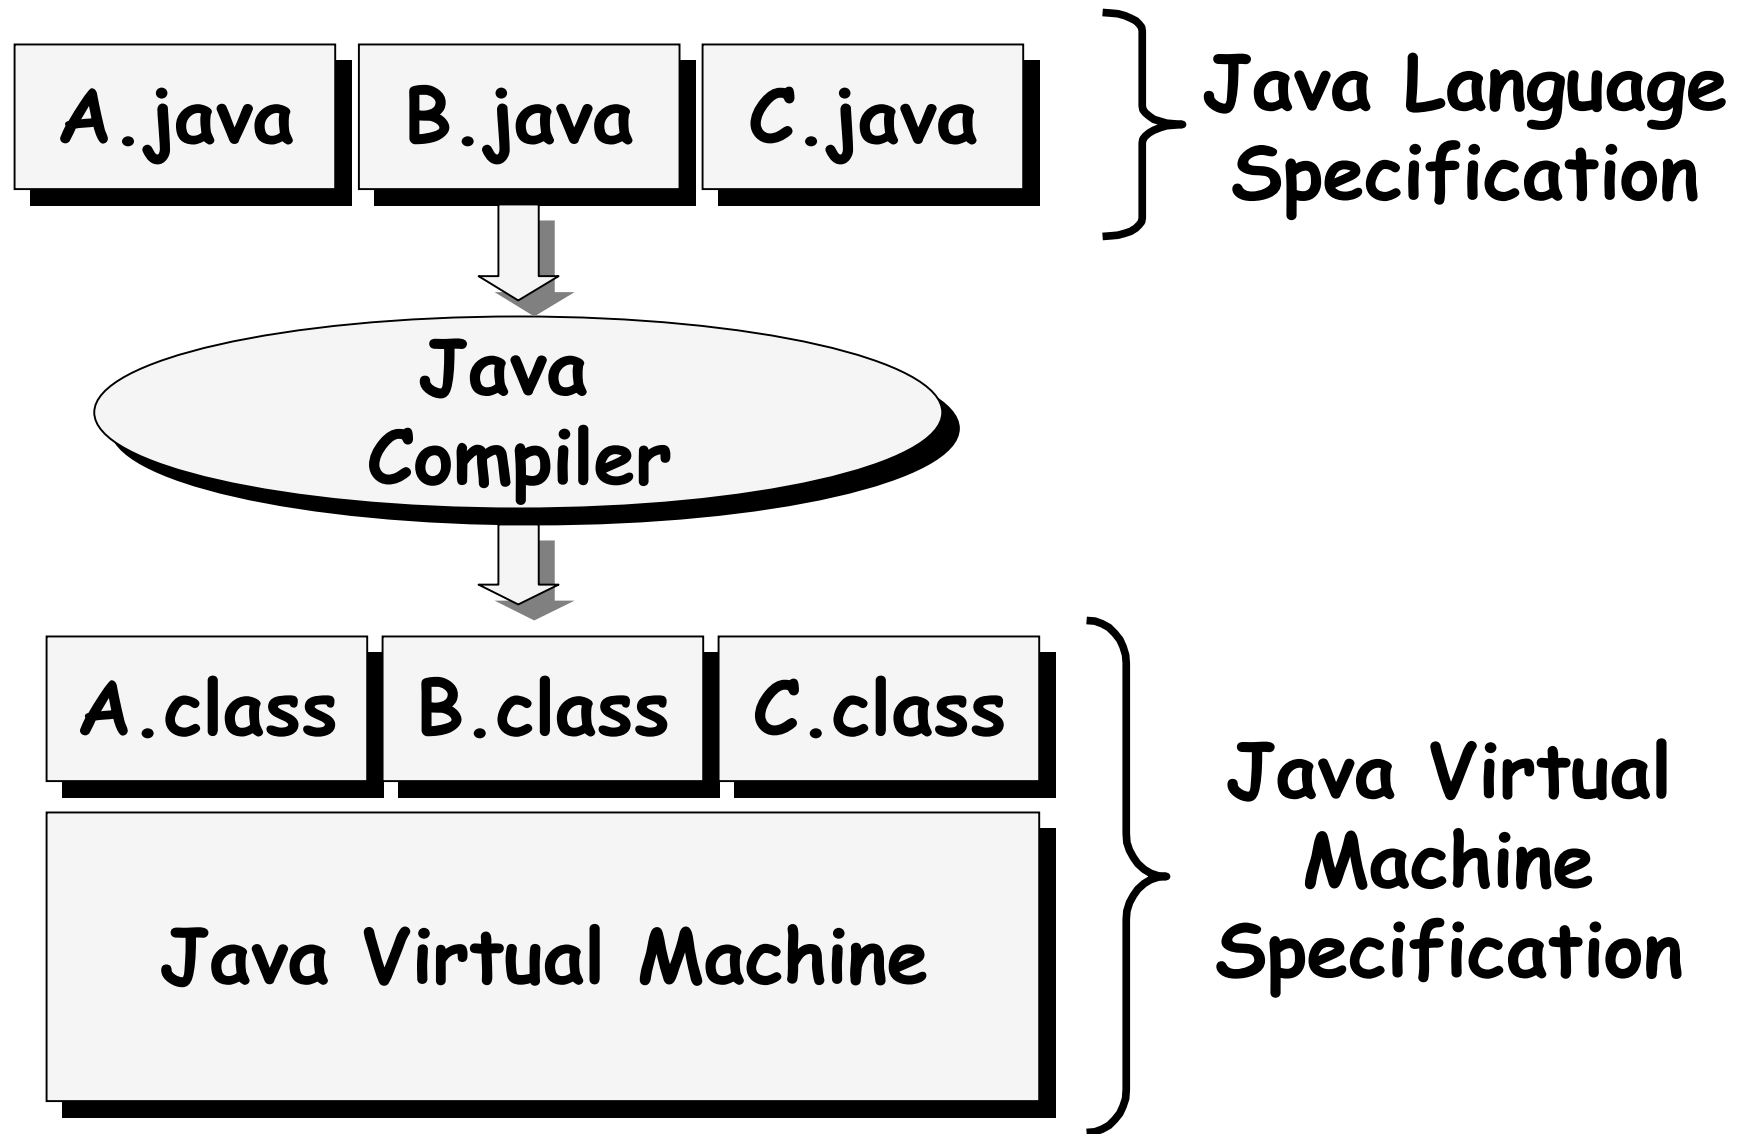


Introduction to Java Virtual Machine

Outline

- Java Language, Java Virtual Machine and Java Platform
- Organization of Java Virtual Machine
- Garbage Collection
- Interpreter and Just-In-Time Compiler

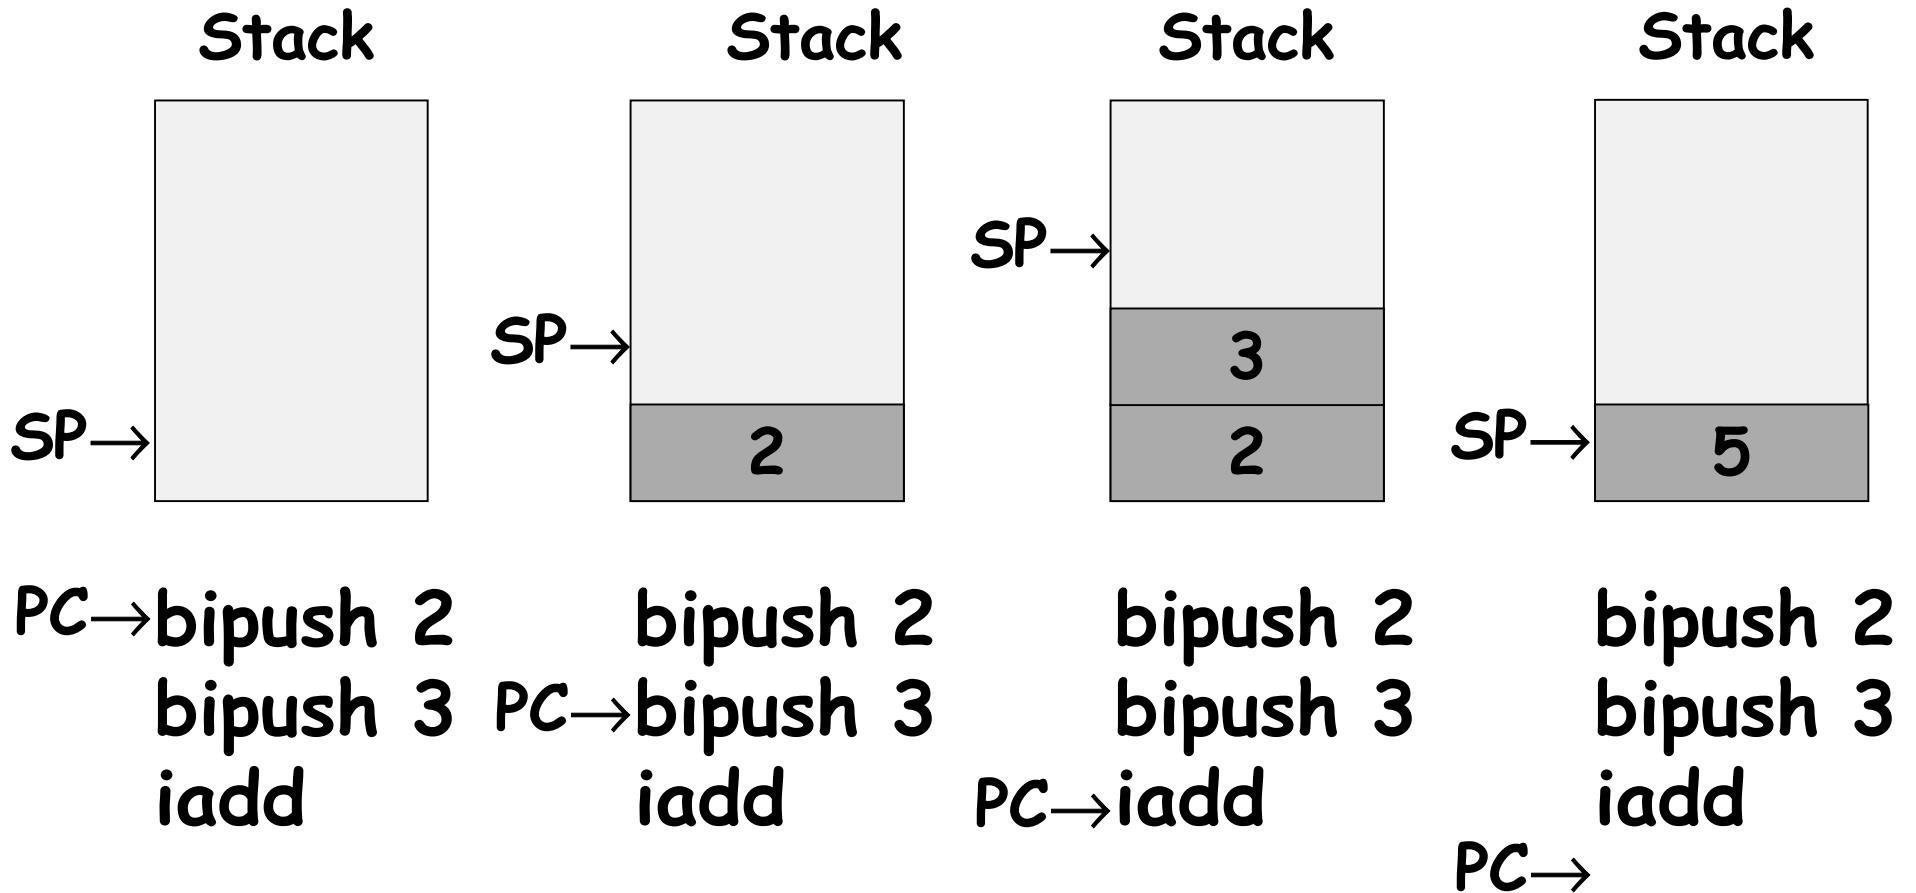
The Big Picture



What Is in the JVM Spec?

- Bytecodes - the instruction set for Java Virtual Machine
- Class File Format - The platform independent representation of Java binary code
- Verification Rules - the algorithm for identifying programs that cannot compromise the integrity of the JVM

Bytecode example



Class File Example (1)

HelloWorld.java

```
public class HelloWorld extends Object {  
    private String s;  
    public HelloWorld() {  
        s = 'Hello World!';  
    }  
    public void sayHello() {  
        System.out.println(s);  
    }  
    public static void main(String[] args) {  
        HelloWorld hello = new HelloWorld();  
        hello.sayHello( );  
    }  
}
```

Class File Example (2)

HelloWorld.class

```
class HelloWorld
  Superclass java/lang/Object
  Constant Pool
    #0: 'Hello World'
  Fields
    s descriptor : Ljava/lang/String;
    modifiers   : private
  Methods
    <init>      descriptor : ()V
                modifiers  : public
    sayHello   descriptor : ()V
                modifiers  : public
    main       descriptor : (Ljava/lang/String[];)V
                modifiers  : public, static
```


Bytecodes

- Bytecodes for <init> (the constructor)
- Bytecodes for sayHello
- Bytecodes for main

Verification

- Is it a structurally valid class file?
- Are all constant references correct?
- Will each instruction always find a correct formed stack and local variable array?
- Check out external references
- Other safety requirements

Implementations of JVM




**JVM Specification does not specify
how a JVM is implemented**

Java Platforms (1)

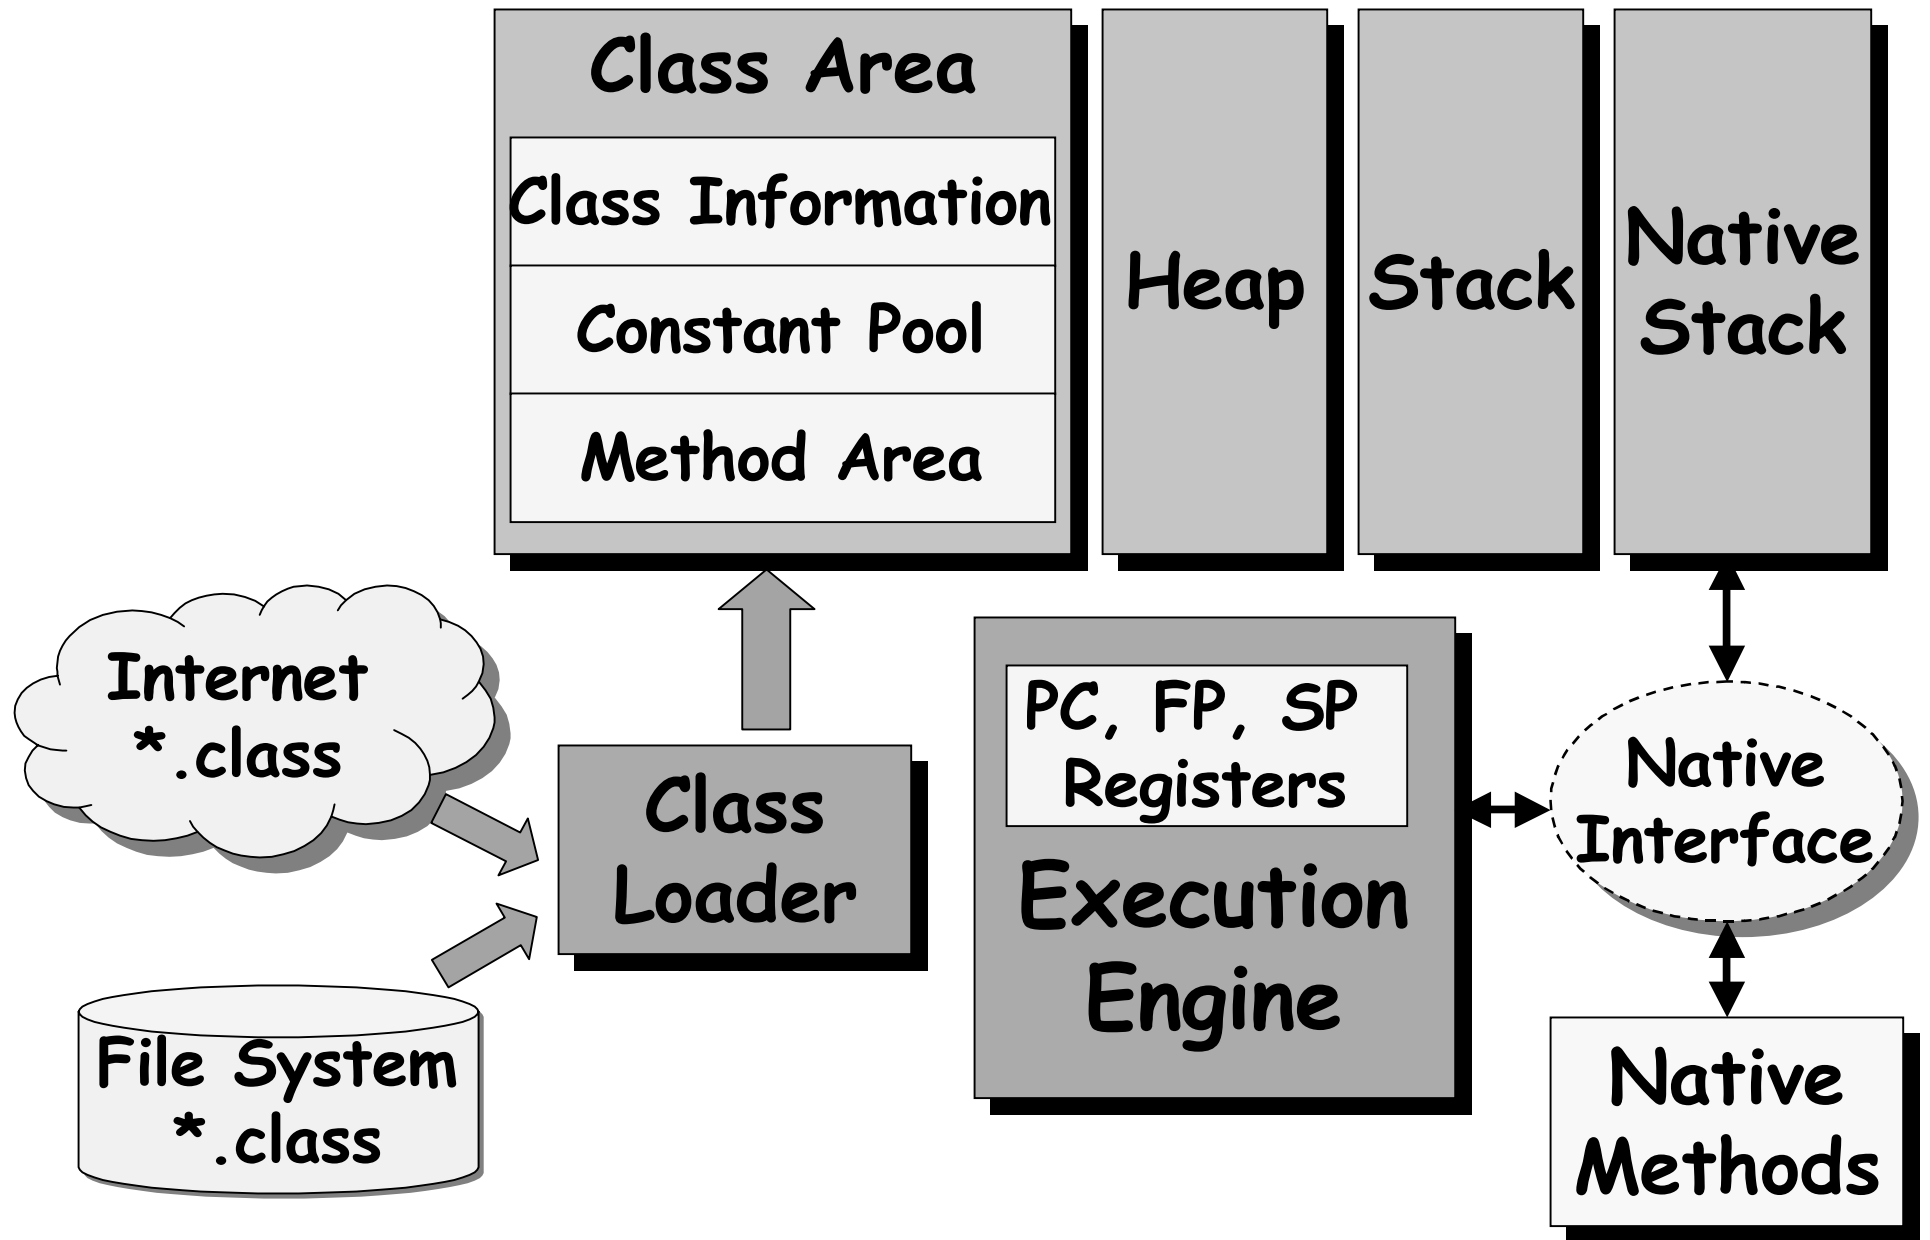
- A Java Platform consists of Java Virtual Machine and a set of standard classes
- JVM in all platforms must satisfy JVM Specification
- Standard classes can be tailored according to the resource constraints
- Three levels of Java platforms: J2EE, J2SE and J2ME

Java Platforms (2)



From KVM White Paper (Sun Microsystems)

Organization of JVM



Execution Engine

- Executes Java bytecodes either using interpreter or Just-In-Time compiler
- Registers:
 - PC: Program Counter
 - FP: Frame Pointer
 - SP: Operand Stack Top Pointer

Class Loader

1. Loading: finding and importing the binary data for a class
2. Linking:
 - Verification: ensuring the correctness of the imported type
 - Preparation: allocating memory for class variables and initializing the memory to default values
 - Resolution: transforming symbolic references from the type into direct references.
3. Initialization: invoking Java code that initializes class variables to their proper starting values

Class Area


- **Class Information**
 - Internal representation of Java classes
 - Information about the superclass and implemented interfaces
 - Information about the fields and methods
- **Constant Pool**
- **Method Area**
 - Contains the bytecodes of the methods of the loaded classes

Stack

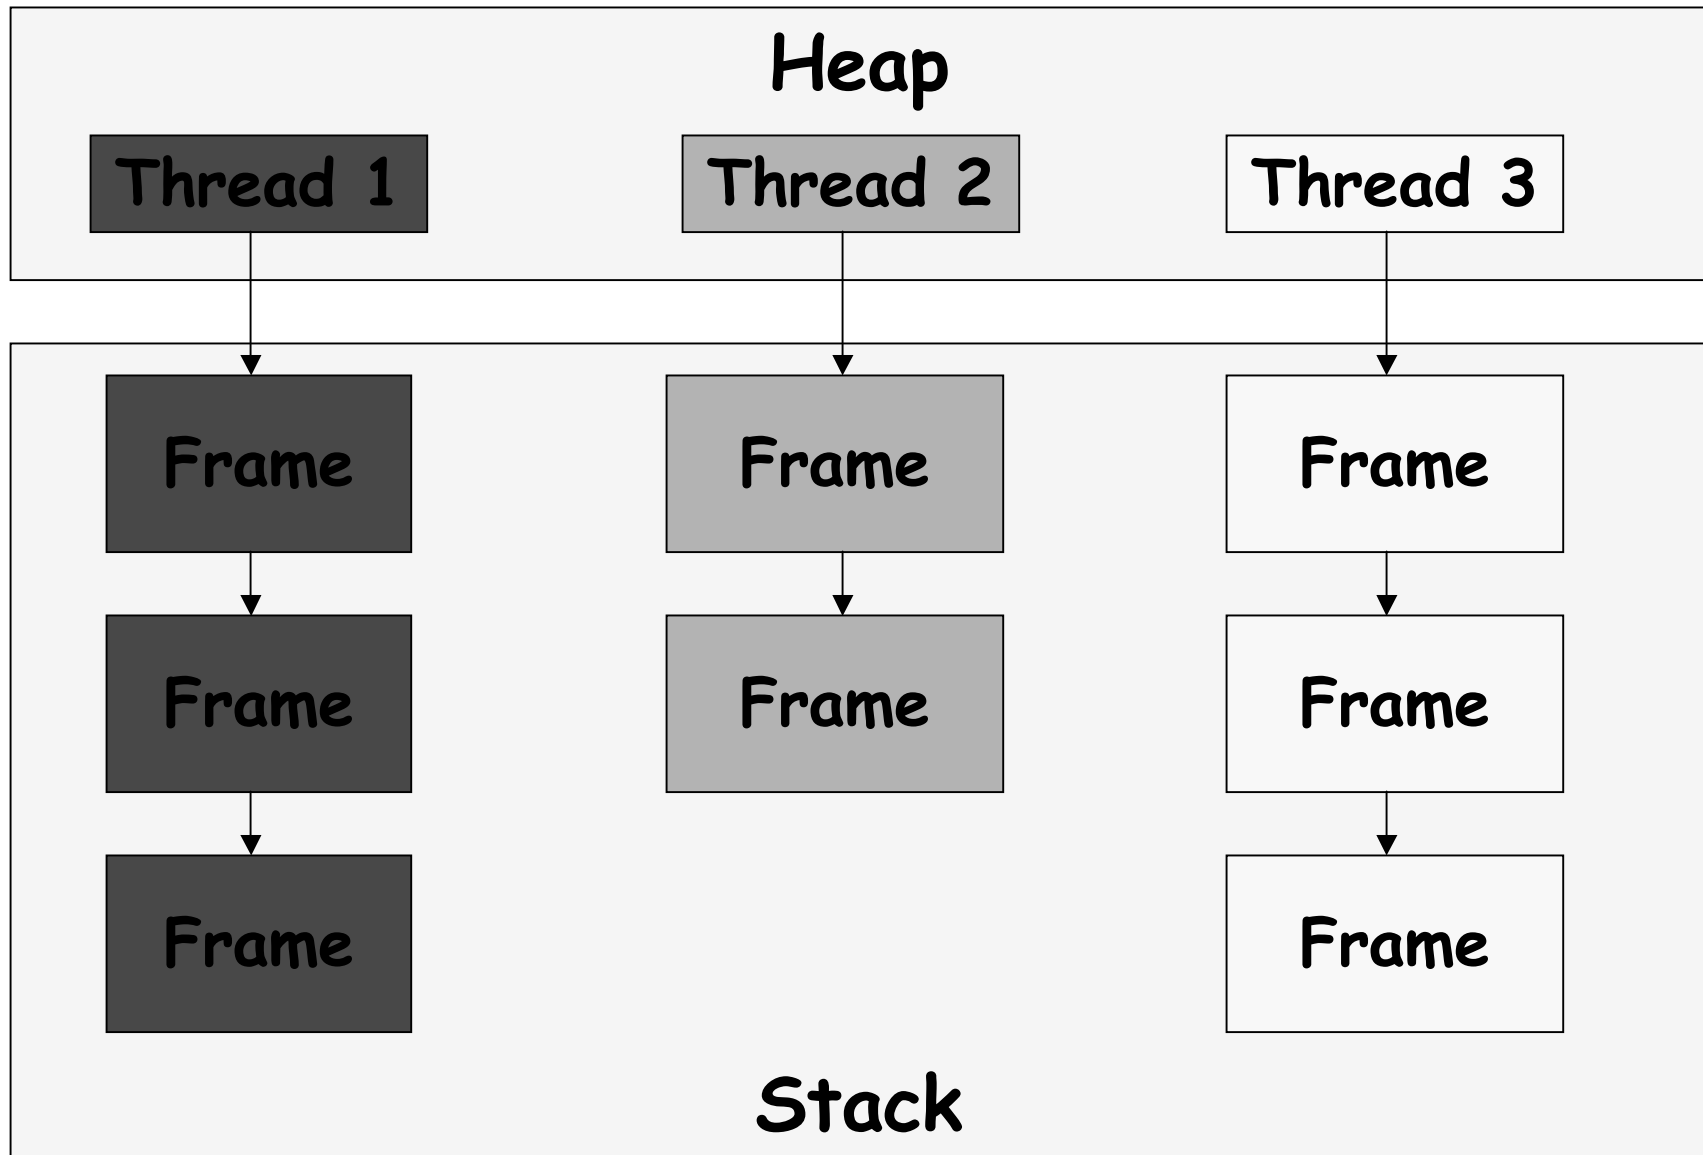
- The Java stack is composed of frames
 - A frame contains the state of one Java method invocation
 - Logically, a frame has two parts: local variable array and operand stack
- JVM has no registers; it uses the operand stack for storage of intermediate data values
 - to keep the JVM's instruction set compact
 - to facilitate implementation on architectures with limited number of general purpose registers
- Each Java thread has its own stack and cannot access the stack of other threads

Stack Frame

```
public class A
{
    ... ..
    void f(int x)
    {
        int i;
        for(i=0; i<x; i++)
        {
            ... ..
        }
        ... ..
    }
}
```




Stack - Each Thread Has its own Stack



Heap

- All Java objects are allocated in the heap
- Java applications cannot explicitly free an object
- The Garbage Collector is invoked from time to time automatically to reclaim the objects that are no longer needed by the application
- The heap is shared by all Java threads

Java Objects in the Heap




Garbage Collector


Roots: internally defined by the JVM implementation

Live Objects: reachable from the roots


Garbage (Dead Objects): not reachable from the roots, not accessible to the application




Mark / Sweep / Compaction




Before GC



After Mark



After Sweep




After Compact

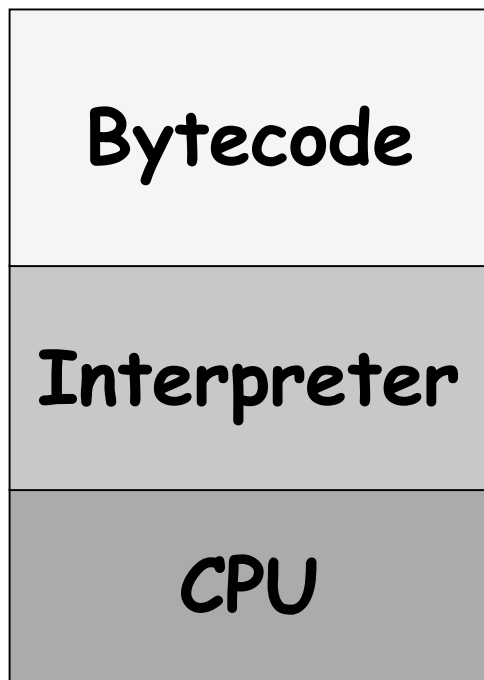
■ Live ■ Garbage □ Unknown ■ Free

Generational Garbage Collection

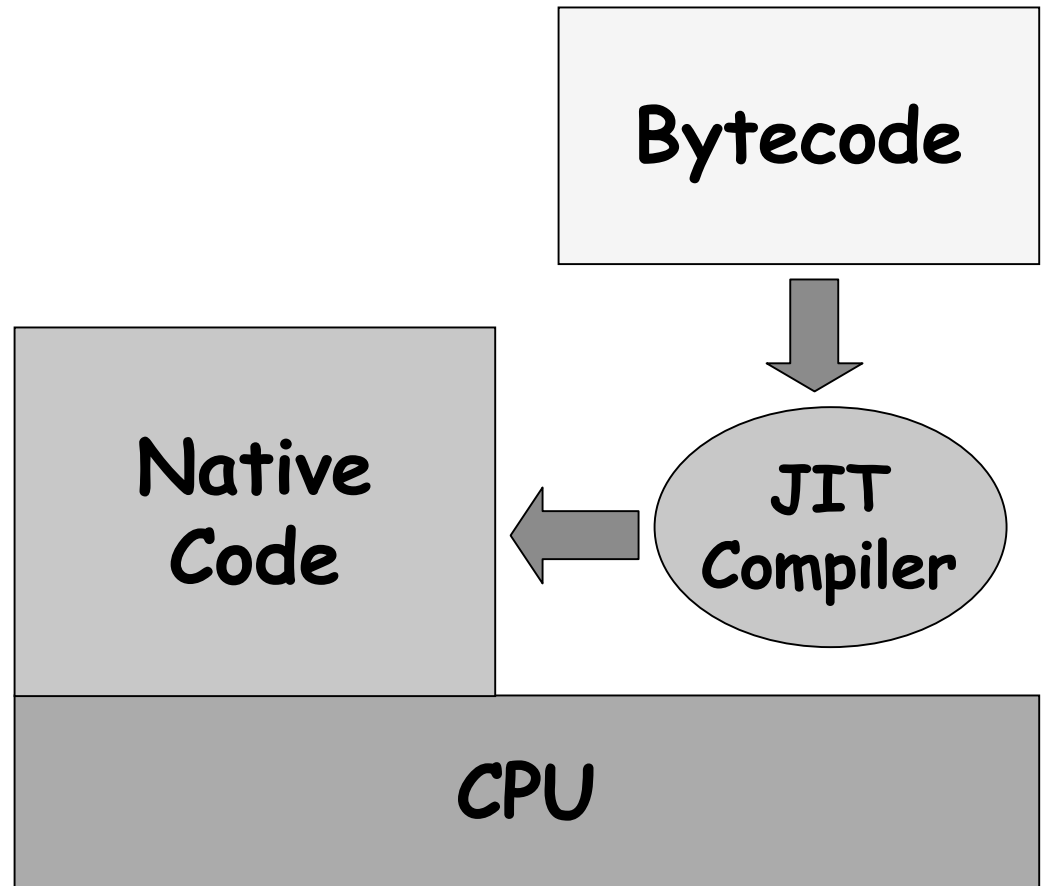
- Most objects live for very short time
- A small percentage of them live much longer



Interpreter vs Just-In-Time Compiler



Interpretation



JIT Compilation

Bytecode Interpreter (1)

```
while(program not end ) {  
    fetch next bytecode => b  
    switch(b) {  
        case ILOAD:  
            load an integer from the local  
            variable array and push on top  
            of current operand stack;  
        case ISTORE:  
            pop an integer from the top of  
            current operand stack and store  
            it into the local variable array;  
        case ALOAD:  
            ... . . .  
    } // end of switch  
} // end of while
```

Bytecode interpreter (2)

- Advantage
 - Ease to implement
 - Does not need extra memory to store compiled code
- Disadvantage
 - Very Slow: 10 ~ 100 times slower than execution of native code

Just-In-Time Compiler

- Dynamically compiles bytecode into native code at runtime, usually in method granularity
- Execution of native code is much faster than interpretation of bytecode
- Compilation is time consuming and may slow down the application
- Tradeoffs between execution time and compilation time

Adaptive Compiler

- Observation: less than 20% of the methods account for more than 80% of execution time
 - Methods contains loop with large number of iteration;
 - Methods that are frequently invoked
- Idea 1: only compile the methods where the application spends a lot of time
- Idea 2: perform advanced compiler optimization for the hottest methods, simple or no compiler optimization for less hot methods

How Adaptive Compiler Works

- Set three thresholds $T1$, $T2$ ($T1 < T2$)
- Each method has a counter that is initialized to 0. Whenever the method is invoked, increase its counter by 1
- The methods with counter lower than $T1$ are executed using interpreter
- When a method's counter reaches $T1$, compile this method with simple optimizations
- When a method's counter reaches $T2$, recompile this method with advanced optimizations