


*Compilation 2012*

# **The Java Virtual Machine**

Jan Midtgaard  
Michael I. Schwartzbach  
Aarhus University


# Virtual Machines in Compilation

---



# Virtual Machines in Compilation

---



# Compiling Virtual Machine Code

---

- Example:
  - gcc translates into RTL, optimizes RTL, and then compiles RTL into native code
- Advantages:
  - facilitates code generators for many targets
- Disadvantage:
  - a code generator must be built for each target

# Interpreting Virtual Machine Code

---

- Examples:
  - P-code for Pascal interpreters
  - Postscript code for display devices
  - Java bytecode for the Java Virtual Machine
- Advantages:
  - easy to generate code
  - the code is architecture independent
  - bytecode can be more compact
- Disadvantage:
  - poor performance (naively 5-100 times slower)

# Designing A Virtual Machine

---

- The instruction set may be more or less high-level
- A balance must be found between:
  - the work of the compiler
  - the work of the interpreter
- In the extreme case, there is only one instruction:
  - compiler guy: `execute "program"`
  - interpreter guy: `print "result"`
- The resulting sweet spot involves:
  - doing as much as possible at compile time
  - exposing the program structure to the interpreter
  - minimizing the size of the generated code
  - being able to verify security&safety policies on compiled code

# Java Virtual Machine

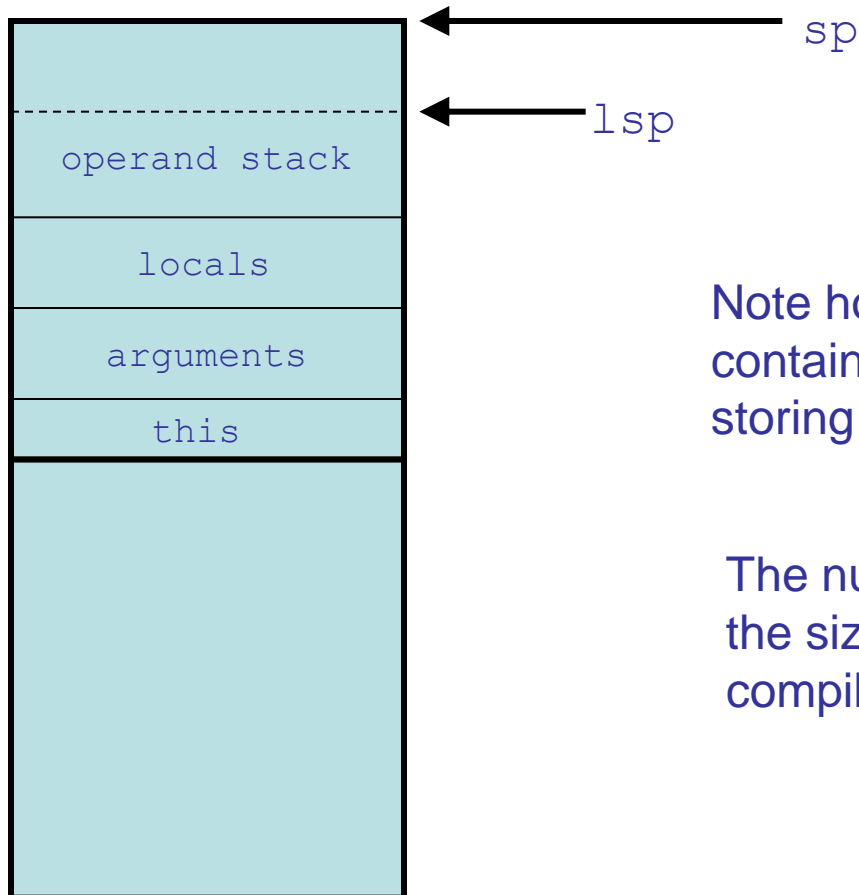
---

- Components of the JVM:
  - stack (per thread)
  - heap
  - constant pool
  - code segment
  - program counter (per thread)

(we ignore multiple threads in this presentation)

# The Java Stack

- The *stack* consists of *frames*:



Note how a frame of the *call stack* contains smaller *operand stack* for storing temporary values


The number of local slots in and the size of a frame are fixed at compile-time



# The Java Heap

---

- The *heap* consists of *objects*:



# The Java Constant Pool


---

- The *constant pool* consists of all *constant data*:
  - numbers
  - strings
  - symbolic names of classes, interfaces, and fields

# The Java Code Segment

---

- The *code segment* consists of *bytecodes* of variable sizes:



# Java Bytecodes

---

- A *bytecode* instruction consists of:
  - a one-byte opcode
  - a variable number of arguments  
(offsets or pointers to the constant pool)
- It consumes and produces some stack elements
- Constants, locals, and stack elements are typed:
  - addresses (*a*)
  - primitive types (*i, c, b, s, f, d, l*)

# Class Files

---

- Java compilers generate class files:
  - magic number (0xCAFEBAFE)
  - minor version/major version
  - constant pool
  - access flags
  - this class
  - super class
  - interfaces
  - fields
  - methods
  - attributes (extra hints for the JVM or other applications)

# Class Loading

---

- Classes are loaded lazily when first accessed
- Class name must match file name
- Super classes are loaded first (transitively)
- The bytecode is verified
- Static fields are allocated and given default values
- Static initializers are executed

# From Methods to Bytecode

- A simple Java method:

```
public int Abs(int i)
{ if (i < 0)
    return(i * -1);
  else
    return(i);
}
```

```
.method public Abs(I)I // int argument, int result
.limit stack 2         // stack with 2 locations
.limit locals 2        // space for 2 locals

                                // --locals--  --stack---
    iload_1                // [ x -3 ]      [ -3  * ]
    ifge Label1            // [ x -3 ]      [  *  * ]
    iload_1                // [ x -3 ]      [ -3  * ]
    iconst_m1              // [ x -3 ]      [ -3 -1 ]
    imul                   // [ x -3 ]      [  3  * ]
    ireturn                // [ x -3 ]      [  *  * ]
Label1:
    iload_1
    ireturn
.end method
```

- Comments show trace of: `x.Abs(-3)`

# A Sketch of a Bytecode Interpreter

The core of a VM consists of a fetch-decode-execute loop:

```
pc = code.start;
while(true)
{
    npc = pc + instruction_length(code[pc]);
    switch (opcode(code[pc]))
    {
        case ILOAD_1: push(locals[1]);
                      break;
        case ILOAD:  push(locals[code[pc+1]]);
                      break;
        case ISTORE: t = pop();
                      locals[code[pc+1]] = t;
                      break;
        case IADD:   t1 = pop(); t2 = pop();
                      push(t1 + t2);
                      break;
        case IFEQ:   t = pop();
                      if (t==0) npc = code[pc+1];
                      break;
        ...
    }
    pc = npc;
}
```



# Instructions

---

- The JVM has 256 instructions for:
  - arithmetic operations
  - branch operations
  - constant loading operations
  - locals operations
  - stack operations
  - class operations
  - method operations
- See the JVM specification for the full list

# Arithmetic Operations

<code>ineg</code>	$[\dots:i] \rightarrow [\dots:-i]$
<code>i2c</code>	$[\dots:i] \rightarrow [\dots:i\%65536]$
<code>iadd</code>	$[\dots:i:j] \rightarrow [\dots:i+j]$
<code>isub</code>	$[\dots:i:j] \rightarrow [\dots:i-j]$
<code>imul</code>	$[\dots:i:j] \rightarrow [\dots:i*j]$
<code>idiv</code>	$[\dots:i:j] \rightarrow [\dots:i/j]$
<code>irem</code>	$[\dots:i:j] \rightarrow [\dots:i\%j]$
<code>iinc k i</code>	$[\dots] \rightarrow [\dots]$ $\text{locals}[k] = \text{locals}[k] + i$

# Branch Operations

goto L	[...] → [...] branch always
--------	--------------------------------

ifeq L	[...:i] → [...] branch if i==0
--------	-----------------------------------

ifne L	[...:i] → [...] branch if i!=0
--------	-----------------------------------

ifnull L	[...:a] → [...] branch if a==null
----------	--------------------------------------

ifnonnull L	[...:a] → [...] branch if a!=null
-------------	--------------------------------------

# Branch Operations

`if_icmpeq L`      `[...:i:j] → [...]`  
branch if `i==j`

`if_icmpne L`

`if_icmpgt L`

`if_icmplt L`

`if_icmpge L`

`if_icmple L`

`if_acmpeq L`      `[...:a:b] → [...]`  
branch if `a==b`

`if_acmpne L`

# Constant Loading Operations

<code>iconst_0</code>	<code>[...] → [...:0]</code>
-----------------------	------------------------------

<code>iconst_1</code>	<code>[...] → [...:1]</code>
-----------------------	------------------------------

...

<code>iconst_5</code>	<code>[...] → [...:5]</code>
-----------------------	------------------------------

<code>aconst_null</code>	<code>[...] → [...:null]</code>
--------------------------	---------------------------------

<code>ldc i</code>	<code>[...] → [...:i]</code>
--------------------	------------------------------

More precisely, the argument of `ldc` is an index into the constant pool of the current class, and the constant at that index is pushed.

<code>ldc s</code>	<code>[...] → [...:String(s)]</code>
--------------------	--------------------------------------

Again, the argument to `ldc` is actually an index into the constant pool

# Locals Operations

<code>iload k</code>	<code>[...] → [...:locals[k]]</code>
----------------------	--------------------------------------

<code>istore k</code>	<code>[...:i] → [...]</code> <code>locals[k]=i</code>
-----------------------	--

<code>aload k</code>	<code>[...] → [...:locals[k]]</code>
----------------------	--------------------------------------

<code>astore k</code>	<code>[...:a] → [...]</code> <code>locals[k]=a</code>
-----------------------	--

# Field Operations

```
getfield f sig [...:a] → [...:a.f]
```

```
putfield f sig [...:a:v] → [...]  
a.f=v
```

```
getstatic f sig [...] → [...:C.f]
```

```
putstatic f sig [...:v] → [...]  
C.f=v
```

More precisely, the argument to these operations is an index in the constant pool which must contain the signature of the corresponding field.

# Stack Operations

dup	$[\dots:v] \rightarrow [\dots:v:v]$
-----	-------------------------------------

pop	$[\dots:v] \rightarrow [\dots]$
-----	---------------------------------

swap	$[\dots:v:w] \rightarrow [\dots:w:v]$
------	---------------------------------------

nop	$[\dots] \rightarrow [\dots]$
-----	-------------------------------

dup_x1	$[\dots:v:w] \rightarrow [\dots:w:v:w]$
--------	---

dup_x2	$[\dots:u:v:w] \rightarrow [\dots:w:u:v:w]$
--------	---



# Class Operations

`new C`                      `[...] → [...:a]`

`instanceof C`      `[...:a] → [...:i]`  
                    `if (a==null) i==0`  
                    `else i==(type(a) ≤ C)`

`checkcast C`      `[...:a] → [...:a]`  
                    `if (a!=null) && !type(a) ≤ C)`  
                    `throw ClassCastException`

# Method Operations

```
invokevirtual name sig
                [...:a:v1:...:vn] → [...(:v)]
m=lookup(type(a),name,sig)
push frame of size m.locals+m.stacksize
locals[0]=a
locals[1]=v1
...
locals[n]=vn
pc=m.code
```

invokestatic

invokespecial

invokeinterface

# Method Operations

---

<code>ireturn</code>	<code>[...:i] → [...]</code> return <code>i</code> and pop stack frame
<code>areturn</code>	<code>[...:a] → [...]</code> return <code>a</code> and pop stack frame
<code>return</code>	<code>[...] → [...]</code> pop stack frame

# A Java Method

---

```
public boolean member(Object item)
{ if (first.equals(item))
    return true;
  else if (rest == null)
    return false;
  else
    return rest.member(item);
}
```

# Generated Bytecode

```
.method public member(Ljava/lang/Object;)Z
.limit locals 2          // locals[0] = x
                          // locals[1] = item
.limit stack 2           // initial stack [ * * ]
aload_0                  // [ x * ]
getfield Cons/first Ljava/lang/Object;
                          // [ x.first *]
aload_1                  // [ x.first item]
invokevirtual java/lang/Object/equals(Ljava/lang/Object;)Z
                          // [bool *]
ifeq else_1              // [ * * ]
iconst_1                 // [ 1 * ]
ireturn                  // [ * * ]
else_1:
aload_0                  // [ x * ]
getfield Cons/rest LCons; // [ x.rest * ]
aconst_null              // [ x.rest null]
if_acmpne else_2         // [ * * ]
iconst_0                 // [ 0 * ]
ireturn                  // [ * * ]
else_2:
aload_0                  // [ x * ]
getfield Cons/rest LCons; // [ x.rest * ]
aload_1                  // [ x.rest item ]
invokevirtual Cons/member(Ljava/lang/Object;)Z
                          // [ bool * ]
ireturn                  // [ * * ]
.end method
```

# Bytecode Verification

---

- Bytecode cannot be trusted to be well-behaved
- Before execution, it must be verified
- Verification is performed:
  - at class loading time
  - at runtime
- A Java compiler must generate verifiable code

# Verification: Syntax

---

- The first 4 bytes of a class file must contain the magic number `0xCAFEBAFE`
- The bytecodes must be syntactically correct

# Verification: Constants and Headers

---

- Final classes are not subclassed
- Final methods are not overridden
- Every class except `Object` has a superclass
- All constants are legal
- Field and method references have valid signatures



# Verification: Instructions

---

- Branch targets are within the code segment
- Only legal offsets are referenced
- Constants have appropriate types
- All instructions are complete
- Execution cannot fall off the end of the code

# Verification: Dataflow Analysis and Type Checking

---

- At each program point, the stack always has the same size and types of objects
- No uninitialized locals are referenced
- Methods are invoked with appropriate arguments
- Fields are assigned appropriate values
- All instructions have appropriate types of arguments on the stack and in the locals

# Verification: Gotcha

```
.method public static main([Ljava/lang/String;)V
.throws java/lang/Exception
.limit stack 2
.limit locals 1
ldc -21248564
invokevirtual java/io/InputStream/read() I
return
```

```
java Fake
```

```
Exception in thread "main" java.lang.VerifyError:
(class: Fake, method: main signature: ([Ljava/lang/String;)V)
Expecting to find object/array on stack
```

# Verification: Gotcha Again

---

```
.method public static main([Ljava/lang/String;)V
.throws java/lang/Exception
.limit stack 2
.limit locals 2
iload_1
return
```

```
java Fake
```

```
Exception in thread "main" java.lang.VerifyError:
(class: Fake, method: main signature: ([Ljava/lang/String;)V)
Accessing value from uninitialized register 1
```

# Verification: Gotcha Once More

---

```
ifeq A  
ldc 42  
goto B  
A:  
ldc "fortytwo"  
B:
```

```
java Fake
```

```
Exception in thread "main" java.lang.VerifyError:  
(class: Fake, method: main signature: ([Ljava/lang/String;)V  
Mismatched stack types
```

# Verification: Gonna Getcha Every Time

---

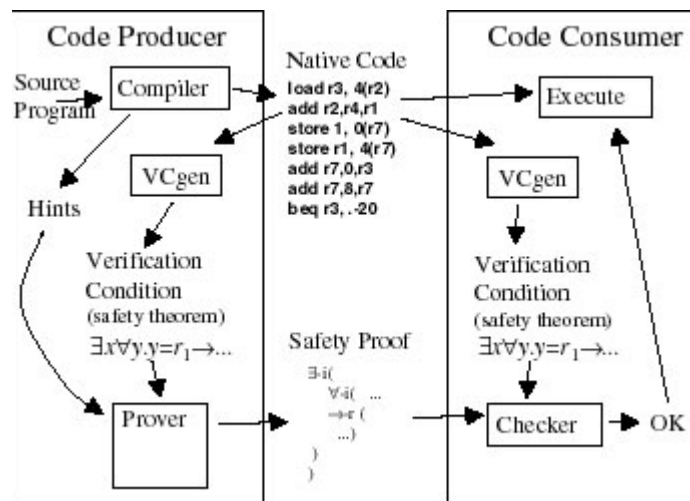
```
A:  
iconst_5  
goto A
```

```
java Fake
```

```
Exception in thread "main" java.lang.VerifyError:  
(class: Fake, method: main signature: ([Ljava/lang/String;)V  
Inconsistent stack height 1 != 0
```


# Alternative: Proof-Carrying Code

- Elegant verification approach to enforce safety and security policies
  - based on theorem proving methods
- E.g., allows distribution of native code while maintaining the safety guarantees of VM code
- No trust in the originator is needed



# JVM Implementation

- A naive bytecode interpreter is slow
- State-of-the-art JVM implementations are not:

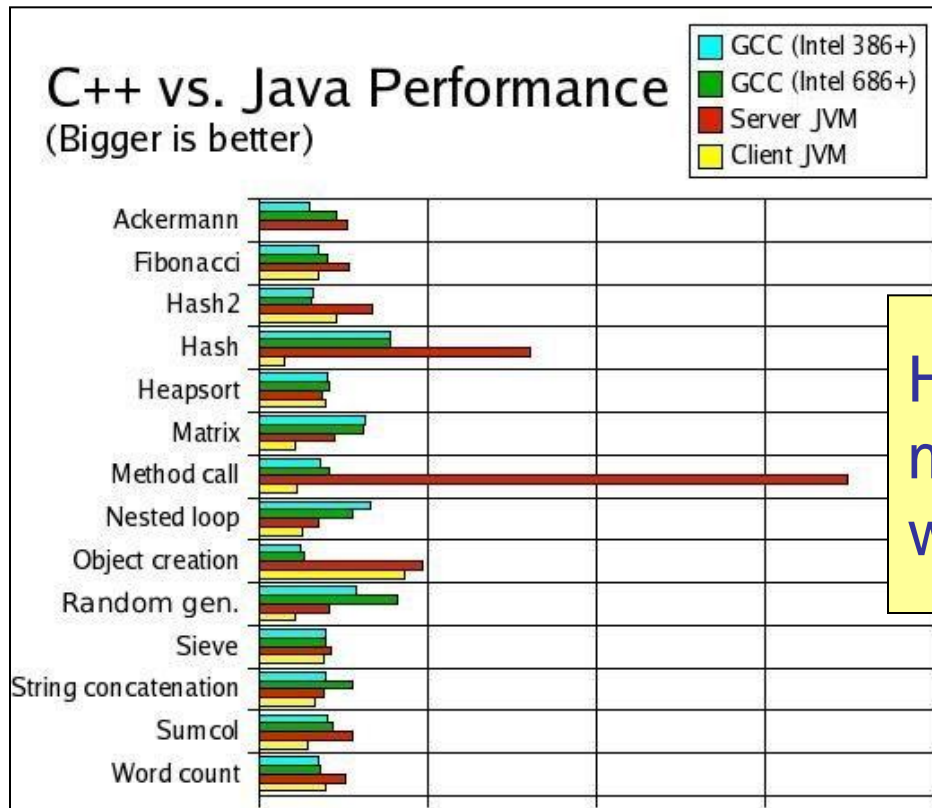


<http://kano.net/javabench>



# JVM Implementation

- A naive bytecode interpreter is slow
- State-of-the-art JVM implementations are not:



However: take  
micro-benchmarks  
with a grain of salt

<http://kano.net/javabench>

# Just-In-Time Compilation

---

- Exemplified by SUN's HotSpot JVM
- Bytecode fragments are compiled at runtime
  - targeted at the native platform
  - based on runtime profiling
  - customization is possible
- Offers more opportunities than a static compiler
- It needs to be fast as it happens at run-time

# Other Java Bytecode Tools

---

- assembler (`jasmin`)
- disassembler (`javap`)
- decompiler (`cavaj`, `mocha`, `jad`)
- obfuscator (dozens of these...)
- analyzer (`soot`)